

Pyodide: scientific Python compiled to WebAssembly

Roman Yurchak



November 11 - 15, 2020



Agenda

Introducing Pyodide

Pyodide UIs and frontend Python apps

Use case: deploying machine learning models

Current challenges and outlook

Introducing Pyodide

What is WebAssembly?

WebAssembly (or WASM) is a binary instruction format for a stack-based virtual machine,

- Initially implemented for browsers
- can also be executed in standalone environments

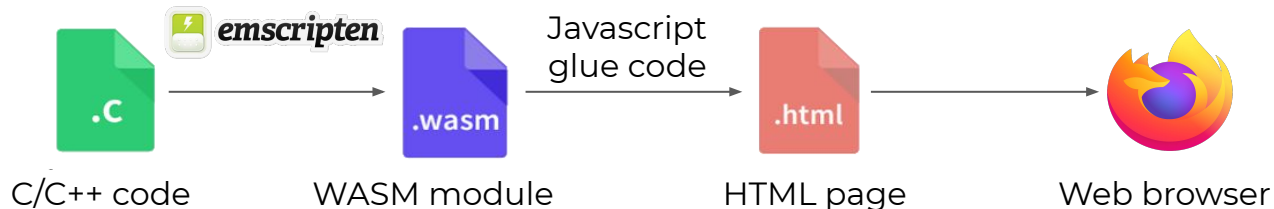
Features

- portable
- near native performance
- sandboxed



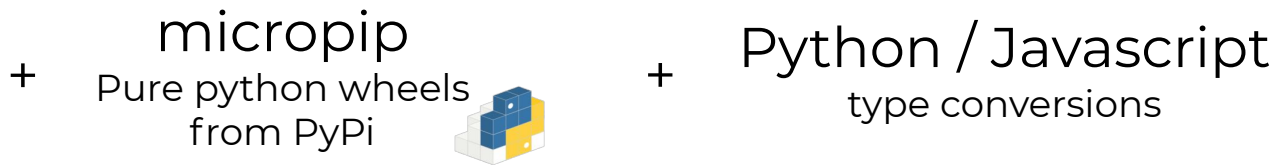
WEBASSEMBLY

Build pipeline



Pyodide project

CPython 3.8 and the scientific Python stack, compiled to WebAssembly



Project repository: github.com/iodide-project/pyodide



Related projects

Several other projects also allow to run Python in the browser ...



pypy.js

PyPy for Python 2 compiled
to asm.js
(No longer maintained)

Brython

Python 3 Javascript implementation
+ parts of stdlib



RustPython

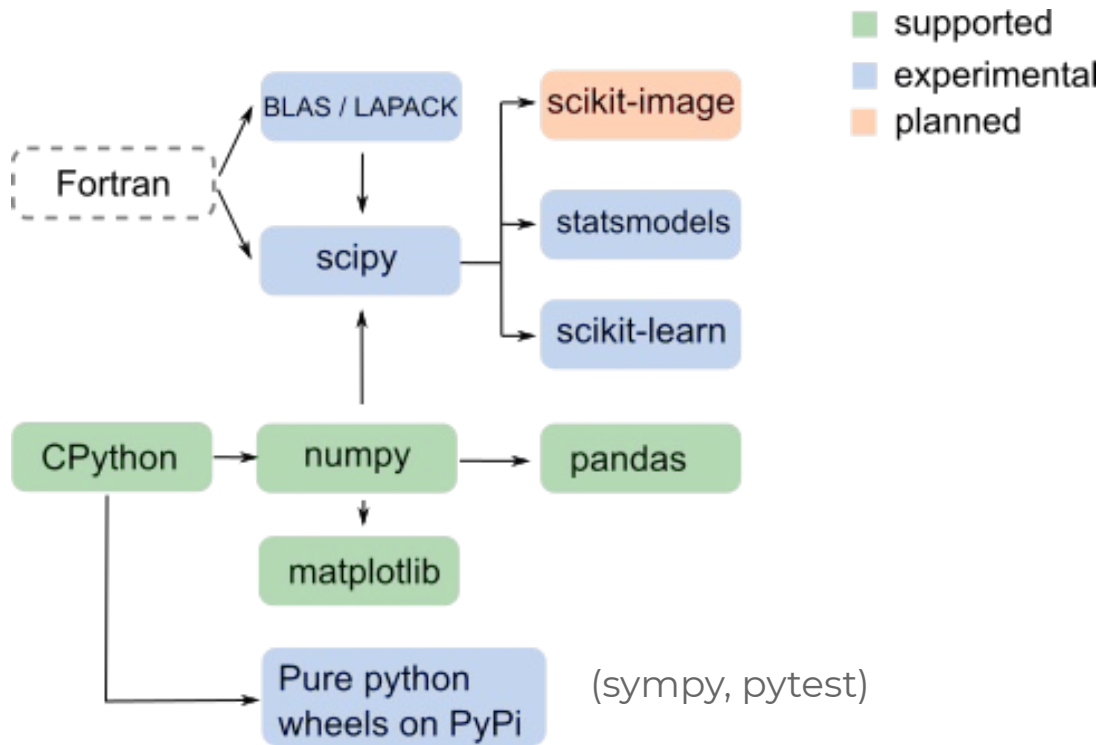
A Python 3 interpreter written in Rust

micropython-ports-wasm

A WASM MicroPython port (experimental)

Supported Python packages in pyodide

that use C extensions



Micropip and Python wheels

Pure Python packages

Installed with **micropip**, if wheels available:

- from PyPi or arbitrary location.
- rudimentary dependency resolution



Examples

See [PEP 427](#):

-py3-none-any.whl -> pure Python wheel

-cp38-manylinux1_x86_64.whl -> Linux wheel
(not compatible with pyodide)

Might still need to use the pyodide build system, to apply patches (e.g. unsupported modules)

Python packages with C extensions

Need to use the pyodide build system (write a **meta.yaml**, similar to conda-forge).

Distributed via pyodide-cdn2.iodide.io (.js & .data files). Can be loaded with `pyodide.loadPackage` (or `micropip`).

Python ↔ Javascript type conversions

Bidirectional type conversions,

- between native types when possible (float, str, list, ..)
- otherwise proxy objects are used with a number of supported operators (getattr, setattr, __call__, ...)

Using Javascript from Python

Allows to access DOM or any JS object in the the global (window) namespace

```
import js
js.document.title = 'New window title'
```

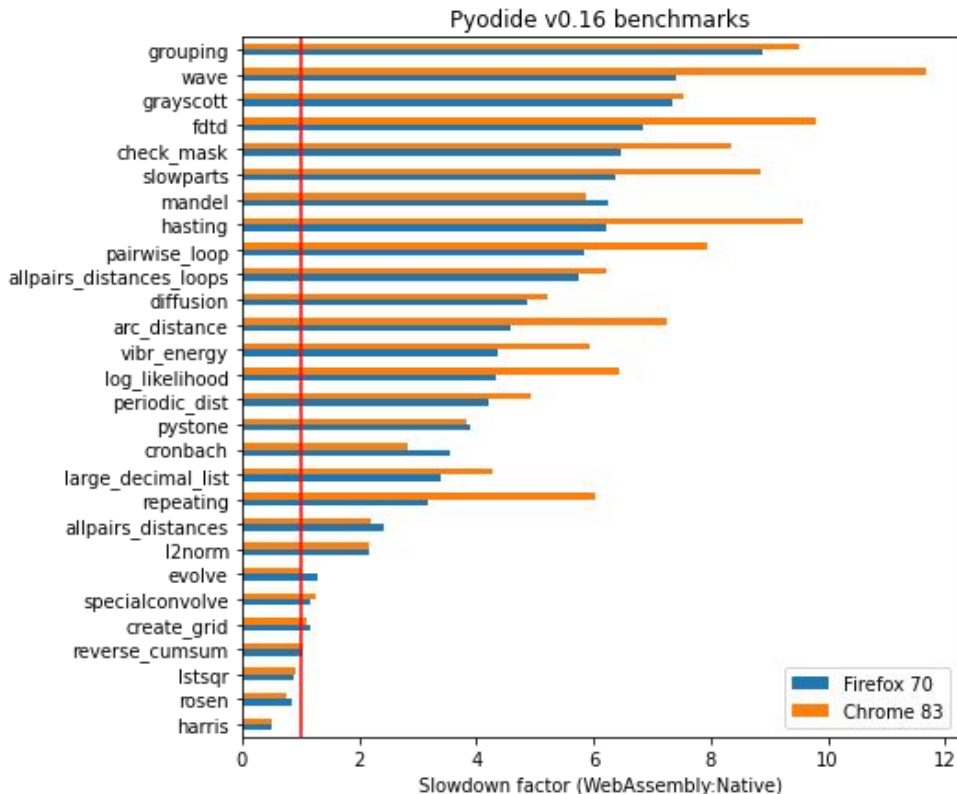
Using Python from Javascript

A Python object (in global scope) can be brought over to Javascript

```
var sys = pyodide.pyimport('sys');
```

Fore more details: pyodide.readthedocs.io/en/latest/type_conversions.html

Performance



Firefox: 4-8 slower for pure Python, 1-2 times slower for C-ext.
Significant speed improvements for WASM in Chrome last year.

Pyodide Uls and frontend Python apps

User Interfaces for pyodide (1)

Pyodide REPL

A simple JS console

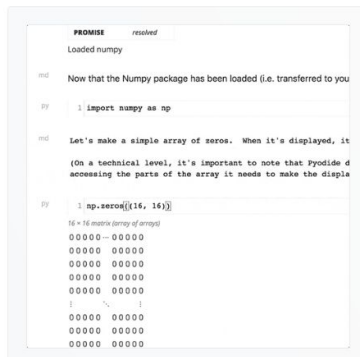
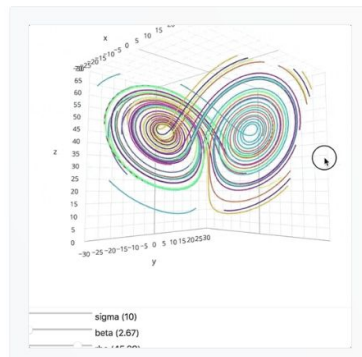
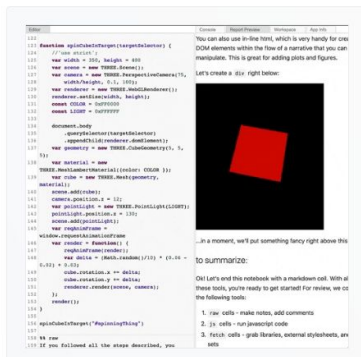
```
Welcome to the Pyodide terminal emulator 🐍
>>> from js import document
>>> document
[object HTMLDocument]
```

Iodide

Literate scientific computing
and communication for the
web

alpha.iodide.io

(no longer actively developed)



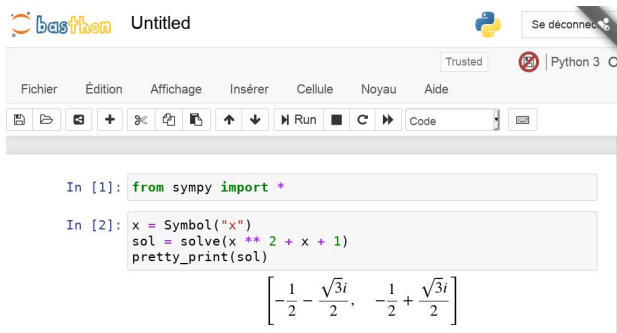
User Interfaces for pyodide (2)

Basthon

Static version of jupyter notebook

notebook.basthon.fr

(first version in French)



Untitled

Fichier Édition Affichage Insérer Cellule Noyau Aide

In [1]: `from sympy import *`

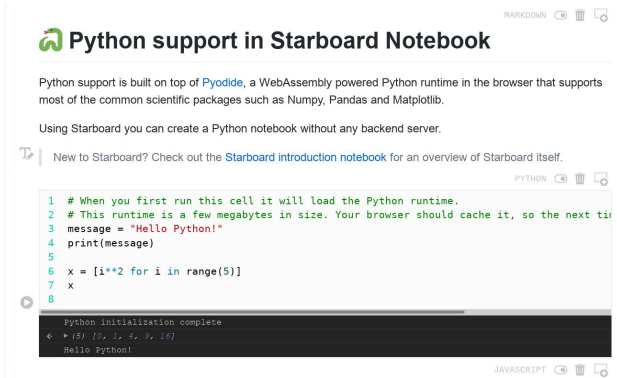
In [2]: `x = Symbol("x")`
`sol = solve(x ** 2 + x + 1)`
`pretty_print(sol)`

$$\left[-\frac{1}{2} - \frac{\sqrt{3}i}{2}, -\frac{1}{2} + \frac{\sqrt{3}i}{2} \right]$$


Starboard Notebook

The shareable in-browser notebook

starboard.gg/#python



Python support in Starboard Notebook

Python support is built on top of [Pyodide](#), a WebAssembly powered Python runtime in the browser that supports most of the common scientific packages such as Numpy, Pandas and Matplotlib.

Using Starboard you can create a Python notebook without any backend server.

New to Starboard? Check out the [Starboard introduction notebook](#) for an overview of Starboard itself.

```
1 # When you first run this cell it will load the Python runtime.
2 # This runtime is a few megabytes in size. Your browser should cache it, so the next ti
3 message = "Hello Python!"
4 print(message)
5
6 x = [i**2 for i in range(5)]
7 x
8
```

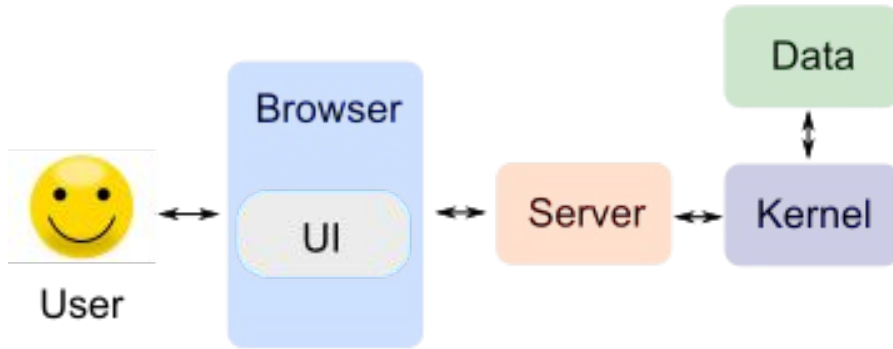
Python initialization complete

↳ (5) [0, 1, 4, 9, 16]
Hello Python!

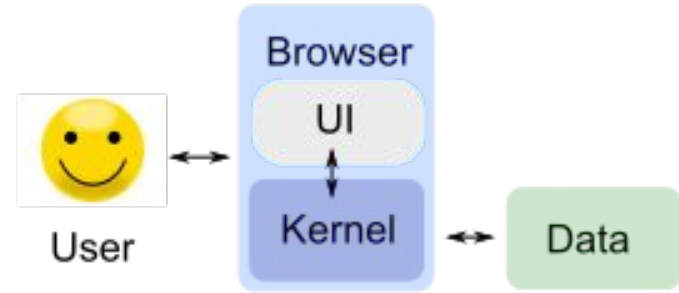


Architecture

Application with a backend server



Application with only static files



Frontend web apps in Python

Usability

Still no Python installation needed, just open a web page

Scalability

Serving static files is easy, scales well to a large number of users

- No need for extensive backend infrastructure / maintenance effort
- *Example:* currently 50k-80k downloads/month for the main pyodide package

Packages only downloaded once, then cached in the browser

Frontend web apps in Python

Privacy

All calculations are run locally, no data needs to be sent to a remote server

Algorithmic Transparency

The app is just an archive with python files/objects

- can be examined on the client side
- can be either an advantage or a disadvantage depending on the use case



Use case: deploying
machine learning models

Deploying machine learning models

Classical workflow

1. Train the machine learning (ML) model
2. Serialize model to disk
3. Develop a web service
4. Package in a container (Docker)
5. Deploy on a server

What format do you use to serialize [@scikit_learn](#) models in production?



133 votes · Final results

Tools for ML inference with WASM support



ONNX



TensorFlow.js



tract

Fast, small model size but restricted to predefined operators...

Deploying scikit-learn models in pyodide*

Use pickle?

Unsafe, brittle to environment changes **but** portable and non opaque

Steps

1. Create an environment with the same Python and dependencies versions as pyodide
2. Pickle the model (`pickle.dumps`) and deserialize it in pyodide (`pickle.loads`)
3. Run inference from JS

Walk through: github.com/rth/notebooks/tree/master/pyodide/PyDataGlobal2020-demo

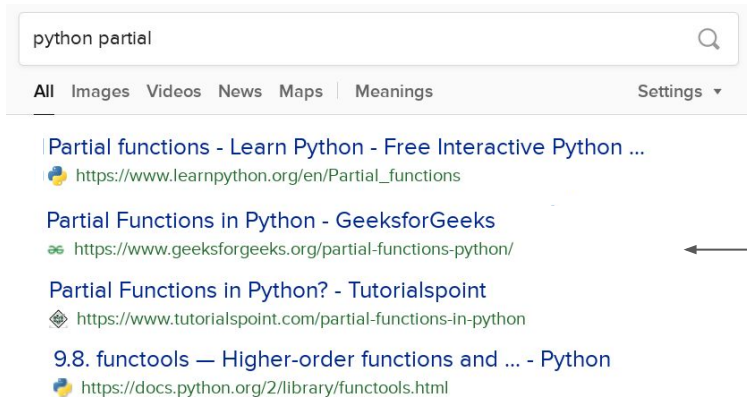
* Very experimental: see next part for the remaining challenges for use in production

Personalized active learning

Not only inference, we could also re-train (or fine tune) models in the browser.

- Smaller, personalized dataset
- Privacy preserving, control over the objective function
- Though, challenging model evaluation

Example



Python 3 stdlib docs,
please!





Current challenges and outlook

Download sizes for pyodide packages

Download size is not an optimisation criterion in the Python ecosystem (unlike for JS)

Historically large packages (e.g. scipy)

- Possibly due the state of Python packaging 10+ years ago

Inclusions of test files in the main package (e.g. `import numpy.tests`)

- Also sometimes test data

Static imports analysis is challenging

- blanket imports in `__init__`, conditional imports

Package	Size	Size (brotli compression)
Python/pyodide core	22 MB	7.6 MB
pandas	16 MB	8.2 MB

Significant optimization potential for the package sizes.

Known limitations

Some of the constraints that require workarounds,

WASM VM

- No subprocess, no threading (WIP)
- No sockets (HTTP requests need to be expressed with JS)
- 32 bit architecture
- Not all syscalls are implemented in emscripten

JavaScript

- No int type, only Number or BigInt
- No standard ND array type
- Can use reserved keywords from Python

`Promise.new(...).then(...).finally(...)` —————> SyntaxError in Python

Outlook

- Keeping up with emscripten releases (many fixes, size and performance improvements)
- Upstreaming patches
- Re-implement some stdlib modules (e.g. http.client) with Web APIs
- Avoid blocking the main JS thread during calculations
- Dynamic linking of LAPACK and better scipy packaging
- Sustainability of the package build system
- Threading (waiting for wider browser adoption)

New contributors are very welcome!

Both technical and non technical.

Team and partners

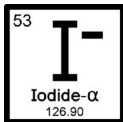
Michael Droettboom

William Lachance

Brendan Colloran

Hamilton Ulmer

Teon Brooks



Jan Max Meyer

Madhur Tandon

Roman Yurchak



Pyodide contributors

Partners and related project



Thank you!

<https://github.com/iodide-project/pyodide>

See PyData Global 2020 program for dates of question sessions.

[@RomanYurchak](#)

roman.yurchak@symerio.com